UC Davis STA 242 2015 Spring Assignment 2

Wenhao Wu, 9987583

April 30, 2015

1 Data Structure and Algorithm Design

A `BMLGrid` instance contains 3 components:

grid A `r`-by-`c` integer matrix. If `grid[i,j]==0`, then there is no car on the crossing of i-th row and j-th column; if `grid[i,j]==1`, there is a red car on that grid; if `grid[i,j]==2`, there is a blue car on that. In our program, when `grid` is indexed, it is treated as a vector (1-D).

blue An integer vector contains the 1-D indices of all blue cars in `grid`.

red An integer vector contains the 1-D indices of all red cars in `grid`.

We define 2 key functions that returns a vector of 1-D indices in `grid`

idx_right() Given an input vector of 1-D indices in `grid`, return a vector of 1-D indices in `grid` for grids to the *right* of the input grids.

idx_up() Given an input vector of 1-D indices in `grid`, return a vector of 1-D indices in `grid` for grids to the *up* of the input grids.

Upon each step, we use `idx_up()`(`idx_right()`) to check in `grid` whether the grids to the up(right) of the grids represented by `blue`(`red`) is occupied, then update the cars' indices `blue`(`red`) and the grid state `grid` accordingly.

The R script file to test the basic design of our code is **Test.R**.

2 Verifying and Profiling

In order to verify the functionality of our BML simulation, we use package 'animation' to generate a short movie showing how cars move on the grid. This function has also been incorporated in our final BMLGrid package.

We have also improved the performance of our code using Rprof(), which originally happens at commit e8fcb63. We demonstrate this process by recreating the code for the original algorithm in **TestProfiling_original.R** and compare it with our improved code which is profiled in **TestProfiling.R**. Again we take the total running time over 10 repetitions of constructing a BMLGrid object with $r = 100$, $c = 99$, $\rho = 0.3$ and the same number of red and blue cars, then run the BML simulation for 10000 steps. The generated profiling files for the original and the modified code are **ProfBMLGridOriginal.out** and **ProfBMLGrid.out**, respectively. When examining the profiling result for the original code, the first few lines are shown in Table 1. We notice that function `ifelse()` has a very high overhead, and confirm this observation by using `lines=''show''` option in `summaryRprof()` function, which shows that `ifelse()` in line 68 and 78 have a self.time of 7.82 and 5.96, respectively. The total sampling.time is 31.98.

To reduce this overhead, we notice that in our design the order of `red` and `blue` vector component of a BMLGrid instance does not matter at all. Consequently, a more effective method to update these two components is to simply concatenate a sub-vector of indices of the cars moved and a sub-vector of indices of the cars unmoved. As a result, our modified code has reduced the sampling.time to 20.9, a good 1/3 of

Table 1: Rpof summary for the original code (by.self).

	self.time	self.pct	total.time	total.pct
"ifelse"	13.46	42.09	13.74	42.96
"eval"	8.82	27.58	31.98	100.00
"%%"	2.88	9.01	2.88	9.01
...	...	...	...	...

improvement in running speed.

(a) $\rho = 0.2$ (b) $\rho = 0.33$

(c) $\rho = 0.38$ (d) $\rho = 0.43$

(e) $\rho = 0.5$

Figure 1: Final state of a 100×99 grid with equal number of blue and red cars after 10000 steps for different car density ρ.

3 Simulation Results

3.1 Behavior of the BML model

We pick a $r = 100$, $c = 99$ grid in which the number of blue cars and red cars are the same. After $N = 10000$ steps, we observe a phase transition in the final state of the grid at $\rho \approx 0.38$. The final states of the grid for $\rho = 0.2, 0.33, 0.38, 0.43, 0.5$ are plotted in Fig. 1. When ρ is small, the grid eventually enters an ordered state where the red and blue cars form a few sparse top-left-to-bottom-right diagonal strips that moves right/upward with little interference. When ρ is large, the grid eventually enters a stale state where the red and blue cars form a few dense top-right-to-bottom-left diagonal strips that stopped moving any more. At $\rho \approx 0.38$, we observed a state transition where after $N = 10000$ steps, the grid is neither in a fully ordered, interference-free mode nor in a complete grid lock.

We further justify this phase transition by looking at the evolution of the grid. In Fig. 2, we demonstrate how the average velocity of blue cars change over time for different ρ. As we can see, when ρ is below a threshold, the smaller ρ is, the quicker the average speed reaches 1, indicating that their is no blockage for blue cars in the grid. When ρ is above a threshold, the larger ρ is, the quicker the average speed drops to 0, indicating there is a total grid lock. When $\rho \approx 0.38$, the average speed varies drastically and takes very long (or maybe infinite) time to reach either the ordered state or the grid lock state. In summary, the 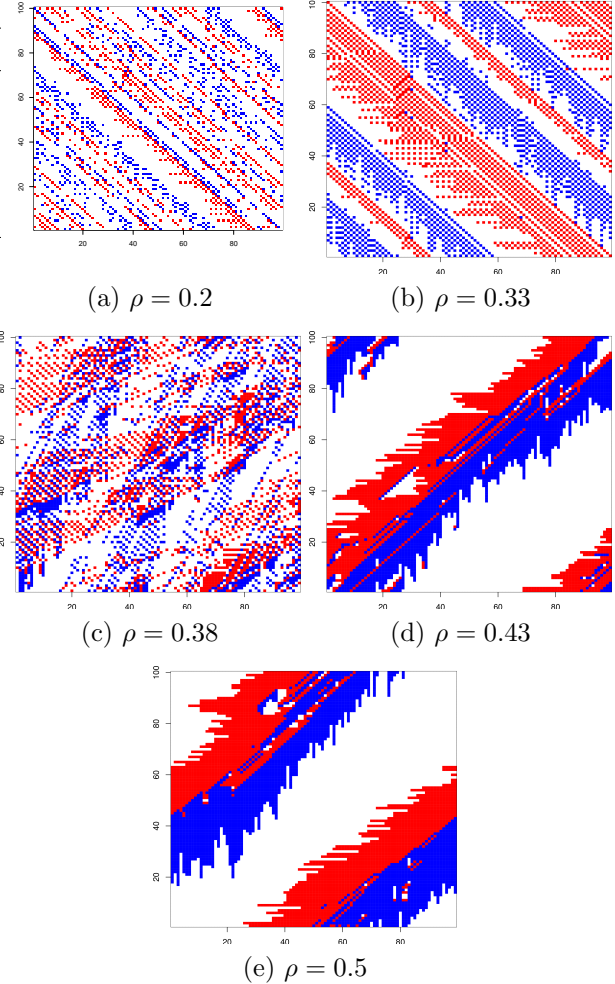BML traffic model demonstrate a *chaotic* behavior: the system is mostly determinstic (except for the grid initialization), however a minor change in parameter ρ will result in a completely different future.

The R script file to demonstrate the behavior of BML model and measure the average velocity are **Test-Behavior.R** and **TestAverageVelocity.R**, respectively.

Figure 2: The evolution of average velocity of the blue cars.

Figure 3: Average user time of a barebone BMLGrid simulation over 5 repetitions for $\rho = 0.1, 0.2, 9.3, 0.4, 0.5, 0.6, 0.7$ and $r = c = 128, 256, 512, 1024$.

3.2 Code Performance

In order to demonstrate the performance of our code, we measure the running time of a function in which

- A BMLGrid instance **g** with $r = c = 128, 256, 512, 1024$ and $\rho = 0.1, 0.2, 9.3, 0.4, 0.5, 0.6, 0.7$ is contructed.

- A BML simulation for $N = 10000$ steps is performed on **g**.

For each of the $4 \times 7 = 28$ instances, the running time is measured by taking the average of 5 runs. We run this test on a Dell Precision T1700 workstation equipped with 16GB RAM and a Core i7-4790K CPU in Ubuntu 14.04 OS. The result is plotted in Fig 3. When ρ is smaller than the threshold, the larger ρ is, the longer the running time is. This is due to the fact that the vector indexing/updating has complexity proportional to the number of cars. When ρ is greater than the threshold, the running time dereases dramatically, thanks to our codes' capability to detect a grid lock state and break from the iteration. As expected, the running time reaches its peak when $\rho \approx 0.38$. Also it is intuitively correct that the running time is about proportional to the size of the grid (i.e. the square of the edge length).

The R script file to test the running time of our BMLGrid key functions is **TestRunningTime.R**.

4 Build BMLGrid Package

The BMLGrid package [1] is developed in RStudio. We manually edit the DESCRIPTION file and the package level object documentation file BMLGrid-package.Rd, while NAMESPACE and other function documentation files are generated with roxygen2. The source package is located in directory **BML-Grid/**.

References

[1] Wenhao Wu. STA 242 Assignment 2: R Package 'BMLGrid'. `https://bitbucket.org/shasqua/stat242_2015_assignment2/src`, 2015. [Online; accessed 30-April-2015].